

TASK-3 (CUSTOMER CHURN PREDICTION)

PROGRAM:

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier,
GradientBoostingClassifier
from sklearn.metrics import accuracy_score, precision_score,
recall_score, f1_score, confusion_matrix, classification_report

# Step 1: Load the dataset
file_path = r"/content/Churn_Modelling.csv" # Update this path
data = pd.read_csv(file_path)

# Inspect the dataset to understand its structure
print(data.head())
print(data.columns)

# Step 2: Data Preprocessing
# Drop unnecessary columns
data = data.drop(columns=['RowNumber', 'CustomerId', 'Surname'])

# Encode categorical variables
label_encoder = LabelEncoder()
data['Gender'] = label_encoder.fit_transform(data['Gender']) # 0
# for Female, 1 for Male
data['Geography'] = label_encoder.fit_transform(data['Geography'])
# Encode Geography

# Separate features and target variable
```

```

X = data.drop(columns=['Exited']) # Features (all columns except
'Exited')
y = data['Exited'] # Target variable (churn: 1 if customer exited,
0 otherwise)

# Standardize the feature values
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Step 3: Split the dataset into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y,
test_size=0.2, random_state=42, stratify=y)

# Step 4: Train and Evaluate Models

# Logistic Regression
log_reg = LogisticRegression(max_iter=1000)
log_reg.fit(X_train, y_train)
y_pred_log_reg = log_reg.predict(X_test)

# Random Forest Classifier
rf = RandomForestClassifier(random_state=42, n_estimators=100)
rf.fit(X_train, y_train)
y_pred_rf = rf.predict(X_test)

# Gradient Boosting Classifier
gb = GradientBoostingClassifier(random_state=42, n_estimators=100)
gb.fit(X_train, y_train)
y_pred_gb = gb.predict(X_test)

# Step 5: Evaluate the models
def evaluate_model(y_test, y_pred, model_name):
    print(f"--- {model_name} ---")
    print(f"Accuracy: {accuracy_score(y_test, y_pred):.4f}")
    print(f"Precision: {precision_score(y_test, y_pred):.4f}")
    print(f"Recall: {recall_score(y_test, y_pred):.4f}")

```

```

print(f"F1 Score: {f1_score(y_test, y_pred):.4f}")
print("\nConfusion Matrix:")
print(confusion_matrix(y_test, y_pred))
print("\nClassification Report:")
print(classification_report(y_test, y_pred))
print("\n")

# Evaluate Logistic Regression
evaluate_model(y_test, y_pred_log_reg, "Logistic Regression")

# Evaluate Random Forest
evaluate_model(y_test, y_pred_rf, "Random Forest")

# Evaluate Gradient Boosting
evaluate_model(y_test, y_pred_gb, "Gradient Boosting")

```

OUTPUT:

RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age
0	1	15634602	Hargrave	619	France	Female
42						
1	2	15647311	Hill	608	Spain	Female
41						
2	3	15619304	Onio	502	France	Female
42						
3	4	15701354	Boni	699	France	Female
39						
4	5	15737888	Mitchell	850	Spain	Female
43						

	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	\
0	2	0.00	1	1		1
1	1	83807.86	1	0		1
2	8	159660.80	3	1		0
3	1	0.00	2	0		0
4	2	125510.82	1	1		1

	EstimatedSalary	Exited
0	101348.88	1
1	112542.58	0
2	113931.57	1
3	93826.63	0
4	79084.10	0

```

Index(['RowNumber', 'CustomerId', 'Surname', 'CreditScore',
      'Geography',
      'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts',
      'HasCrCard',
      'IsActiveMember', 'EstimatedSalary', 'Exited'],
      dtype='object')

```

--- Logistic Regression ---

Accuracy: 0.8050
Precision: 0.5859
Recall: 0.1425
F1 Score: 0.2292

Confusion Matrix:

```

[[1552  41]
 [ 349  58]]

```

Classification Report:

	precision	recall	f1-score	support
0	0.82	0.97	0.89	1593
1	0.59	0.14	0.23	407
accuracy			0.81	2000
macro avg	0.70	0.56	0.56	2000
weighted avg	0.77	0.81	0.75	2000

--- Random Forest ---

Accuracy: 0.8640
Precision: 0.7848
Recall: 0.4570
F1 Score: 0.5776

Confusion Matrix:

```

[[1542  51]

```

[221 186]]

Classification Report:

	precision	recall	f1-score	support
0	0.87	0.97	0.92	1593
1	0.78	0.46	0.58	407
accuracy			0.86	2000
macro avg	0.83	0.71	0.75	2000
weighted avg	0.86	0.86	0.85	2000

--- Gradient Boosting ---

Accuracy: 0.8675
Precision: 0.7886
Recall: 0.4767
F1 Score: 0.5942

Confusion Matrix:

[[1541 52]
[213 194]]

Classification Report:

	precision	recall	f1-score	support
0	0.88	0.97	0.92	1593
1	0.79	0.48	0.59	407
accuracy			0.87	2000
macro avg	0.83	0.72	0.76	2000
weighted avg	0.86	0.87	0.85	2000